

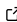


calzone: A Python package for measuring calibration of probabilistic models for classification

Kwok Lung Fan¹, Gene Pennello¹, Qi Liu¹, Nicholas Petrick¹, Ravi K. Samala¹, Frank W. Samuelson¹, Yee Lam Elim Thompson¹, and Qian Cao^{1¶}

¹ U.S. Food and Drug Administration ¶ Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Open Journals](#) 

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#))

Summary

calzone is a Python package for evaluating the calibration of probabilistic outputs of classifier models. It provides a set of functions and classes for visualizing calibration and computing calibration metrics given a representative dataset with the model's predictions and true class labels. The metrics provided in calzone include: Expected Calibration Error (ECE), Maximum Calibration Error (MCE), Hosmer-Lemeshow (HL) statistic, Integrated Calibration Index (ICI), Spiegelhalter's Z-statistics and Cox's calibration slope/intercept. The package is designed with versatility in mind. For many of the metrics, users can adjust the binning scheme and toggle between top-class or class-wise calculations.

Statement of need

Classification is one of the most fundamental tasks in machine learning. Classification models are often evaluated by a proper scoring rule, such as the cross-entropy or mean square error. Examination of the distinguishing power (resolution), such as AUC or Se/Sp are also used to evaluate the model performance. However, the reliability or calibration performance of the model is often overlooked.

Bröcker (2009) has shown that the proper scoring rule can be decomposed into the resolution and reliability. That means even if the model has high resolution (high AUC), it may not be a reliable or calibrated model. In many high-risk machine learning applications, such as medical diagnosis, the reliability of the model is of paramount importance.

We refer calibration as the agreement between the predicted probability and the true posterior probability of a class-of-interest, $P(D = 1|\hat{p} = p) = p$. This is defined as moderate calibration by Calster & Steyerberg (2018) .

In the calzone package, we provide a set of functions and classes for calibration visualization and metrics computation. Existing libraries such as scikit-learn are often not dedicated to calibration metrics computation and don't provide calibration metrics computation that are widely used in the statistical literature. Most libraries for calibration are focusing on calibrating the model instead of measuring the level of calibration with various metrics. calzone is dedicated to calibration metrics computation and visualization.

35 **Functionality**

36 **Reliability Diagram**

37 Reliability Diagram is a graphical representation of the calibration of a classification model
38 (Bröcker & Smith, 2007). It groups the predicted probabilities into bins and plots the mean
39 predicted probability against the empirical frequency in each bin. The reliability diagram can
40 be used to assess the calibration of the model and to identify any systematic errors in the
41 predictions. In addition, we add the option to plot with error bars to show the confidence
42 interval of the empirical frequency in each bin. The error bars are calculated using Wilson's
43 score interval (Wilson, 1927). We provide an example simulated dataset in the example_data
44 folder using beta-binomial distribution (Griffiths, 1973). Users can generate simulated data
45 using the fake_binary_data_generator class in the utils module.

```
from calzone.utils import reliability_diagram
from calzone.vis import plot_reliability_diagram

wellcal_dataloader = data_loader(
    data_path="example_data/simulated_welldata.csv"
)

reliability, confidence, bin_edges, bin_counts = reliability_diagram(
    wellcal_dataloader.labels,
    wellcal_dataloader.probs,
    num_bins=15,
    class_to_plot=1
)

plot_reliability_diagram(
    reliability,
    confidence,
    bin_counts,
    error_bar=True,
    title='Class 1 reliability diagram for well calibrated data'
)
```

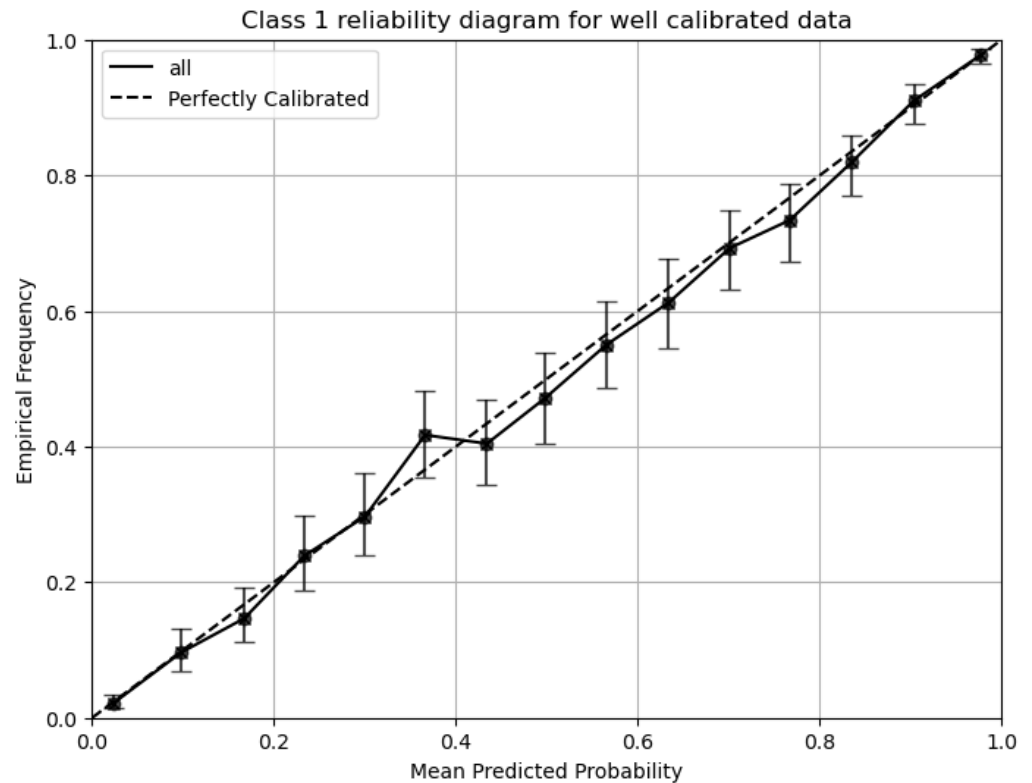


Figure 1: Reliability Diagram for well calibrated data

Calibration metrics

calzone provides functions to compute various calibration metrics. calzone also has a CalibrationMetrics() class which allows the user to compute the calibration metrics in a more convenient way. The following are the metrics that are currently supported in calzone:

Expected Calibration Error (ECE) and Maximum Calibration Error (MCE)

Expected Calibration Error (ECE), Maximum Calibration Error (MCE) and binning-based methods (Guo et al., 2017; Pakdaman Naeini et al., 2015) aim to measure the average deviation between predicted probability and true probability. We provide the option to use equal-width binning or equal-count binning, labeled as ECE-H and ECE-C respectively. Users can also choose to compute the metrics for the class-of-interest or the top-class. In the case of class-of-interest, the program will treat it as a 1-vs-rest classification problem. It can be computed in calzone as follows:

```
from calzone.metrics import calculate_ece_mce

reliability, confidence, bin_edges, bin_counts = reliability_diagram(
    wellcal_dataloader.labels,
    wellcal_dataloader.probs,
    num_bins=10,
    class_to_plot=1,
    is_equal_freq=False
)

ece_h_classone, mce_h_classone = calculate_ece_mce(
```

```

reliability,
confidence,
bin_counts=bin_counts
)

```

58 Hosmer-Lemeshow statistic (HL)

59 Hosmer-Lemeshow statistic (HL) is a statistical test for the calibration of a probabilistic model.
60 It is a chi-square based test that compares the observed and expected number of events in
61 each bin. The null hypothesis is that the model is well calibrated. HL-test first bins data into
62 predicted probability bins (equal-width H or equal-count C) and the test statistic is calculated
63 as:

$$HL = \sum_{m=1}^M \frac{(O_{1,m} - E_{1,m})^2}{E_{1,m}(1 - \frac{E_{1,m}}{N_m})} \sim \chi_{M-2}^2$$

64 where $E_{1,m}$ is the expected number of class-of-interest events in the m^{th} bin, $O_{1,m}$ is the
65 observed number of class-of-interest events in the m^{th} bin, N_m is the total number of
66 observations in the m^{th} bin, and M is the number of bins. In calzone, the HL-test can be
67 computed as follows:

```

from calzone.metrics import hosmer_lemeshow_test

HL_H_ts, HL_H_p, df = hosmer_lemeshow_test(
    reliability,
    confidence,
    bin_count=bin_counts
)

```

68 When performing the HL test on validation sets that are not used in training, the degree of
69 freedom of the HL test changes from $M - 2$ to M . Intuitively, $\frac{(O_{1,m} - E_{1,m})^2}{E_{1,m}(1 - \frac{E_{1,m}}{N_m})}$ is the difference
70 squared divided by the variance of a binomial distribution and follows a chi-square distribution
71 with 1 degree of freedom. Hence, the sum of M chi-square distributions with 1 degree of
72 freedom is a chi-square distribution with M degrees of freedom if the data has no effect on
73 the model. In calzone, user can sepecify the degree of freedom of the HL test by setting the
74 df parameter.

75 Cox's calibration slope/intercept

76 Cox's calibration slope/intercept is a non-parametric method for assessing the calibration of a
77 probabilistic model (COX, 1958). A new logistic regression model is fitted to the data, with
78 the predicted odds ($\frac{p}{1-p}$) as the dependent variable and the true probability as the independent
79 variable. The slope and intercept of the regression line are then used to assess the calibration
80 of the model. A slope of 1 and intercept of 0 indicates perfect calibration. To test whether
81 the model is calibrated, fix the slope to 1 and fit the intercept. If the intercept is significantly
82 different from 0, the model is not calibrated. Then, fix the intercept to 0 and fit the slope.
83 If the slope is significantly different from 1, the model is not calibrated. In calzone, Cox's
84 calibration slope/intercept can be computed as follows:

```

from calzone.metrics import cox_regression_analysis

cox_slope, cox_intercept, cox_slope_ci, cox_intercept_ci = cox_regression_analysis(
    wellcal_data_loader.labels,
    wellcal_data_loader.probs,
    class_to_calculate=1,
    print_results=True,
)

```

```
fix_slope=True
)
```

85 The values of the slope and intercept give you a sense of the form of miscalibration. A slope
86 greater than 1 indicates that the model is overconfident at high probabilities and underconfident
87 at low probabilities, and vice versa. An intercept greater than 0 indicates that the model is
88 overconfident in general, and vice versa. Notice that even if the slope is 1 and the intercept is
89 0, the model might not be calibrated, as Cox's calibration analysis fails to capture some types
90 of miscalibration, including quadratic effects or other non-linearities.

91 Integrated calibration index (ICI)

92 The integrated calibration index (ICI) is very similar to Expected calibration error (ECE). It
93 also tries to measure the average deviation between predicted probability and true probability.
94 However, ICI does not use binning to estimate the true probability of a group of samples with
95 similar predicted probability. Instead, ICI uses curve smoothing techniques to fit the regression
96 curve and uses the regression result as the true probability (Austin & Steyerberg, 2019). The
97 ICI is then calculated using the following formula:

$$ICI = \frac{1}{n} \sum_{i=1}^n |f(p_i) - p_i|$$

98 where f is the fitting function and p is the predicted probability. The curve fitting is usually
99 done with loess regression. However, it is possible to use any curve fitting method to calculate
100 the ICI. In calzone, we provide Cox's ICI and loess ICI support while the user can also use any
101 curve fitting method to calculate the ICI using functions in calzone.

```
from calzone.metrics import (
    cox_regression_analysis,
    lowess_regression_analysis,
    cal_ICI_cox
)

### calculating cox ICI
cox_ici = cal_ICI_cox(
    cox_slope,
    cox_intercept,
    wellcal_data_loader.probs,
    class_to_calculate=1
)

### calculating loess ICI
loess_ici, lowess_fit_p, lowess_fit_p_correct = lowess_regression_analysis(
    wellcal_data_loader.labels,
    wellcal_data_loader.probs,
    class_to_calculate=1,
    span=0.5,
    delta=0.001,
    it=0
)
```

102 Notice that flexible curve fitting methods such as loess regression are very sensitive to the
103 choice of span and delta parameters. The user can visualize the fitting result to avoid overfitting
104 or underfitting.

105 Spiegelhalter's Z-test

106 Spiegelhalter's Z-test is a test of calibration proposed by Spiegelhalter in 1986 (Spiegelhalter,
107 1986). It uses the fact that the Brier score can be decomposed into:

$$B = \frac{1}{N} \sum_{i=1}^N (x_i - p_i)^2 = \frac{1}{N} \sum_{i=1}^N (x_i - p_i)(1 - 2p_i) + \frac{1}{N} \sum_{i=1}^N p_i(1 - p_i)$$

108 And the TS of Z test is defined as:

$$Z = \frac{B - E(B)}{\sqrt{\text{Var}(B)}} = \frac{\sum_{i=1}^N (x_i - p_i)(1 - 2p_i)}{\sum_{i=1}^N (1 - 2p_i)^2 p_i (1 - p_i)}$$

109 and it is asymptotically distributed as a standard normal distribution. In calzone, it can be
110 calculated using:

```
from calzone.metrics import spiegelhalter_z_test

z, p_value = spiegelhalter_z_test(
    wellcal_dataloader.labels,
    wellcal_dataloader.probs,
    class_to_calculate=1
)
```

111 Metrics class

112 calzone also provides a class called CalibrationMetrics() to calculate all the metrics men-
113 tioned above. The user can also use this class to calculate the metrics.

```
from calzone.metrics import CalibrationMetrics

metrics = CalibrationMetrics(class_to_calculate=1)

CalibrationMetrics.calculate_metrics(
    wellcal_dataloader.labels,
    wellcal_dataloader.probs,
    metrics='all'
)
```

114 Other features

115 Bootstrapping

116 calzone also provides bootstrapping to calculate the confidence intervals of the metrics. The
117 user can specify the number of bootstrap samples and the confidence level.

```
from calzone.metrics import CalibrationMetrics

metrics = CalibrationMetrics(class_to_calculate=1)

CalibrationMetrics.bootstrap(
    wellcal_dataloader.labels,
    wellcal_dataloader.probs,
    metrics='all',
    n_samples=1000
)
```

118 and it will return a structured numpy array.

119 Subgroup analysis

120 calzone will perform subgroup analysis by default in the command line user interface. If the
121 user input CSV file contains a subgroup column, the program will compute metrics for the
122 entire dataset and for each subgroup.

123 Prevalence adjustment

124 calzone also provides prevalence adjustment to account for prevalence changes between
125 training data and testing data. Since calibration is defined using posterior probability, a
126 mere shift in the prevalence of the testing data will result in miscalibration. It can be fixed
127 by searching for the optimal derived original prevalence such that the adjusted probability
128 minimizes a proper scoring rule such as cross-entropy loss. The formula of prevalence adjusted
129 probability is:

$$P'(D = 1|\hat{p} = p) = \frac{\eta'/(1-\eta')}{(1/p-1)(\eta/(1-\eta))} = p'$$

130 where η is the prevalence of the testing data, η' is the prevalence of the training data, and p
131 is the predicted probability (Chen et al., 2018; Gu & Pepe, 2010; Horsch et al., 2008; Tian et
132 al., 2020). We search for the optimal η' that minimizes the cross-entropy loss.

133 Multiclass extension

134 calzone also provides multiclass extension to calculate the metrics for multiclass classification.
135 The user can specify the class to calculate the metrics using a 1-vs-rest approach and test
136 the calibration of each class. Alternatively, the user can transform the data and make the
137 problem become a top-class calibration problem. The top-class calibration has a similar format
138 to binary classification, but the class 0 probability is defined as 1 minus the probability of the
139 class with the highest probability, and the class 1 probability is defined as the probability of
140 the class with the highest probability. The labels are transformed into whether the predicted
141 class equals the true class, 0 if not and 1 if yes. Notice that the interpretation of some metrics
142 may change in the top-class transformation.

143 Command line interface

144 calzone also provides a command line interface to calculate the metrics. The user can visualize
145 the calibration curve, calculate the metrics and their confidence intervals using the command
146 line interface. To use the command line interface, the user can run `python cal_metrics.py`
147 `-h` to see the help message.

148 Acknowledgements

149 The authors acknowledge the Research Participation Program at the Center for Devices and
150 Radiological Health administered by the Oak Ridge Institute for Science and Education through
151 an interagency agreement between the U.S. Department of Energy and the U.S. Food and
152 Drug Administration (FDA).

153 Conflicts of interest

154 The authors declare no conflicts of interest.

155 References

156 Austin, P. C., & Steyerberg, E. W. (2019). The integrated calibration index (ICI) and related
157 metrics for quantifying the calibration of logistic regression models. *Statistics in Medicine*,

- 38(21), 4051–4065. <https://doi.org/https://doi.org/10.1002/sim.8281>
- Bröcker, J. (2009). Reliability, sufficiency, and the decomposition of proper scores. *Quarterly Journal of the Royal Meteorological Society*, 135(643), 1512–1519. <https://doi.org/https://doi.org/10.1002/qj.456>
- Bröcker, J., & Smith, L. A. (2007). Increasing the reliability of reliability diagrams. *Weather and Forecasting*, 22(3), 651–661. <https://doi.org/10.1175/WAF993.1>
- Calster, B. V., & Steyerberg, E. W. (2018). Calibration of prognostic risk scores. In *Wiley StatsRef: Statistics reference online* (pp. 1–10). John Wiley & Sons, Ltd. <https://doi.org/https://doi.org/10.1002/9781118445112.stat08078>
- Chen, W., Sahiner, B., Samuelson, F., Pezeshk, A., & Petrick, N. (2018). Calibration of medical diagnostic classifier scores to the probability of disease. *Statistical Methods in Medical Research*, 27(5), 1394–1409. <https://doi.org/10.1177/0962280216661371>
- COX, D. R. (1958). Two further applications of a model for binary regression. *Biometrika*, 45(3-4), 562–565. <https://doi.org/10.1093/biomet/45.3-4.562>
- Griffiths, D. A. (1973). Maximum likelihood estimation for the beta-binomial distribution and an application to the household distribution of the total number of cases of a disease. *Biometrics*, 29(4), 637–648. <http://www.jstor.org/stable/2529131>
- Gu, W., & Pepe, M. S. (2010). Estimating the diagnostic likelihood ratio of a continuous marker. *Biostatistics*, 12(1), 87–101. <https://doi.org/10.1093/biostatistics/kxq045>
- Guo, C., Pleiss, G., Sun, Y., & Weinberger, K. Q. (2017). On calibration of modern neural networks. In D. Precup & Y. W. Teh (Eds.), *Proceedings of the 34th international conference on machine learning* (Vol. 70, pp. 1321–1330). PMLR. <https://proceedings.mlr.press/v70/guo17a.html>
- Horsch, K., Giger, M. L., & Metz, C. E. (2008). Prevalence scaling: Applications to an intelligent workstation for the diagnosis of breast cancer. *Academic Radiology*, 15(11), 1446–1457. <https://doi.org/https://doi.org/10.1016/j.acra.2008.04.022>
- Pakdaman Naeini, M., Cooper, G., & Hauskrecht, M. (2015). Obtaining well calibrated probabilities using bayesian binning. *Proceedings of the AAAI Conference on Artificial Intelligence*, 29(1). <https://doi.org/10.1609/aaai.v29i1.9602>
- Spiegelhalter, D. J. (1986). Probabilistic prediction in patient management and clinical trials. *Statistics in Medicine*, 5(5), 421–433.
- Tian, J., Liu, Y.-C., Glaser, N., Hsu, Y.-C., & Kira, Z. (2020). Posterior re-calibration for imbalanced datasets. In H. Larochelle, M. Ranzato, R. Hadsell, M. F. Balcan, & H. Lin (Eds.), *Advances in neural information processing systems* (Vol. 33, pp. 8101–8113). Curran Associates, Inc. https://proceedings.neurips.cc/paper_files/paper/2020/file/5ca359ab1e9e3b9c478459944a2d9ca5-Paper.pdf
- Wilson, E. B. (1927). Probable inference, the law of succession, and statistical inference. *Journal of the American Statistical Association*, 22(158), 209–212.